

Catch Me If You Can: A Decade of Evasive Malware Attack and Defense

Alexei Bulazel & Bülent Yener

River Loop Security
Rensselaer Polytechnic Institute (RPI)



<https://git.io/vbgEz>

RPISEC

Bio



- Researcher at River Loop Security
- 2015 RPI/RPISEC graduate
 - Work was began as part of my MS and continued until early 2017
 - Collaboration with advisor Dr. Bülent Yener
- Good to be back at ShmooCon!

"A Survey On Automated Dynamic Malware Analysis Evasion and Counter-Evasion: PC, Mobile, and Web" published at ROOTS (Reversing and Offensive-Oriented Trends Symposium) 2017 in Vienna, Austria

Twitter: @0xAlexei

<https://git.io/vbgEz>

Introduction

- Automated dynamic malware analysis is essential to keep up with modern malware (and potentially malicious software)
- **Problem:** malware can detect and evade analysis
- **Solution:** detect or mitigate anti-analysis



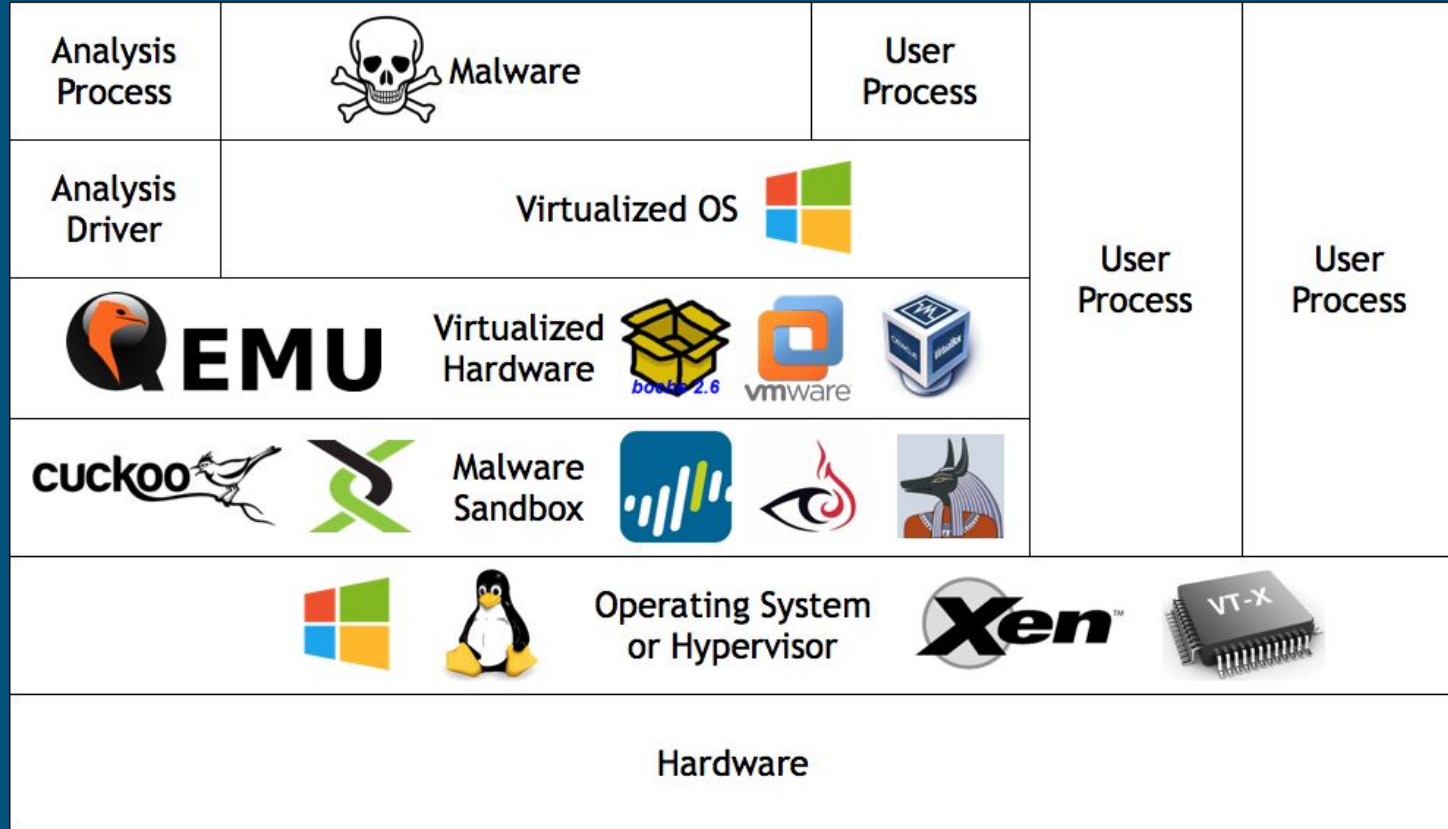
<https://git.io/vbgEz>

Dynamic Automated Analysis Systems

a.k.a:
“malware sandboxes”
“detonation chambers”

Not in scope:

- anti-debug
- anti-RE
- obfuscation



<https://git.io/vbgEz>

Motivation

- Evasive malware techniques is a perennial security conference topic
- Defensive techniques to counter evasion are generally relegated to academic publications or trade secrets
- There's been over a decade of academic work in this field, no one had comprehensively looked at it
- Writing papers is fun

Scope

- Survey of ~200 works on evasive malware techniques, detection, mitigation, and case studies
- Mostly academic works, with a few industry talks and publications
- This presentation - focus on PC-based malware experimentation
- Timeline
 - First papers on evasive malware ~2005
 - Part of a longer lineage, see historical work on software protection



<https://git.io/vbgEz>

Takeaways

- Evasive malware and defenders continually evolve to counter one another
- The fight between malware and analysis systems is likely to continue long into the future
- There are immense challenges to experimental evaluation and the ability to establish ground truth

Read The Paper!

<https://git.io/vbgEz>

[https://github.com/bulaza/
Publications/blob/master/ROOTS2017](https://github.com/bulaza/Publications/blob/master/ROOTS2017)

<https://git.io/vbgEz>

Presentation Outline

1. Introduction
2. **Offense - Detecting Analysis Systems**
3. Defense - Detecting Malware Evasion
4. Defense - Mitigating Malware Evasion
5. Discussion
6. Conclusion



Offense - Detecting Analysis Systems

- Fingerprint Classes
 - Environmental Artifacts
 - Timing
 - CPU Virtualization
 - Process Introspection
 - Reverse Turing Tests
 - Network Artifacts
 - Mobile Sensors
 - Browser Specific

```
bool beingAnalyzed = DetectAnalysis();  
  
if (beingAnalyzed)  
{  
    BehaveBenignly();  
}  
else  
{  
    InfectSystem();  
}
```

Environmental Artifacts & Timing



- Unique distinguishing characteristics of the analysis environment itself
 - Usernames
 - System settings
 - Date
 - Installed software
 - Files on disk
 - Running processes
 - Number of CPUs
 - Amount of RAM
- Timing discrepancies in analysis systems
- Sources:
 - Emulation / virtualization overhead
 - Analysis instrumentation overhead
 - Overhead of physical hardware instrumentation (potentially)
- Challenging to mitigate
 - Garfinkle et al: “extreme engineering hardship and huge runtime overhead”

CPU Virtualization & Process Introspection

- CPU “Red Pills”
- Discrepancies in CPU behavior introduced by virtualization
 - Erroneously accepted/rejected instructions
 - Incorrect exception behavior
 - Flag edge cases
 - MSRs
 - CPUID/SIDT/SGDT/etc discrepancy
- Particularly applicable for emulators
- Discrepancies in internal state
 - Memory or register contents
 - Function hooks
 - Injected libraries
 - Page permission eccentricities
- Commonly used in anti-DBI



Reverse Turing Tests & Network Artifacts

- *Computer* decides if *it* is interacting with computer or human
- Passive: mouse movement, typing cadence, process churn, scrolling
- Active: user must click a button
- Wear-and-Tear: evidence of human use, copy-paste clipboard, “recently opened” file lists, web history, phone camera photos
- Fixed IP address
- Network isolation
- Incorrectly emulated network devices or protocols
- Unusually fast internet service



Detection - Discussion

- Variety of sources: underlying technologies facilitating analysis, system configuration, analysis instrumentation
- Easy to use = easy to mitigate
- Difficult to use = difficult to mitigate
- Reverse Turing Tests seem to be growing in relevance, and are extremely difficult to mitigate against

<https://git.io/vbgEz>

Presentation Outline

1. Introduction
2. Offense - Detecting Analysis Systems
3. Defense - Detecting Malware Evasion
4. Defense - Mitigating Malware Evasion
5. Discussion
6. Conclusion



Detecting Malware Evasion

- Detecting that malware exhibits evasive behavior under dynamic analysis, but not mitigating evasion
 - Comparatively fewer works relative to mitigation work
- Early work - detecting known anti-analysis-techniques
 - 2008: Lau et al.'s DSD-Tracer
- Most works use multi-system execution
 - Run malware in multiple systems and compare behavior offline - discrepancies may indicate evasion in one or more of these systems

Multi-System Execution

- Instruction-level (2009: Kang et al.)
 - Too low level, prone to detect spurious differences
- System call-level (2010: Balzarotti et al. / 2015: Kirat & Vigna - MalGene)
 - Higher level than just instructions
 - MalGene uses algorithms taken from bioinformatics work in protein alignment
- Persistent changes to system state (2011: Lindorfer et al. - Disarm)
 - Jaccard distance-based comparisons
- Behavioral profiling (2014: Kirat et al. - BareCloud)
 - *What* malware did vs. *how* it did it, “hierarchical similarity” algorithms from computer vision and text similarity research

Evasion Detection - Discussion

- Evolution over time to increasingly complex algorithmic approaches, working over increasingly abstracted execution traces
- Multi-system execution is a common solution for evasion detection
- Offline algorithms do not detect evasion in real time
- Detection does not solve the main challenge of evasion, so there is less work in the field compared to mitigation research

<https://git.io/vbgEz>

Presentation Outline

1. Introduction
2. Offense - Detecting Analysis Systems
3. Defense - Detecting Malware Evasion
4. Defense - Mitigating Malware Evasion
5. Discussion
6. Conclusion



Defense - Mitigating Evasion

- Mitigating evasive behavior in malware so that analysis can proceed unhindered
- Early approaches
 - Binary Modification
 - Hiding Environment Artifacts
 - State Modification
 - Multi-Platform Record and Replay
- Path Exploration
- Hypervisor-based Analysis
- Bare Metal Analysis & SMM-based Analysis
- Discussion

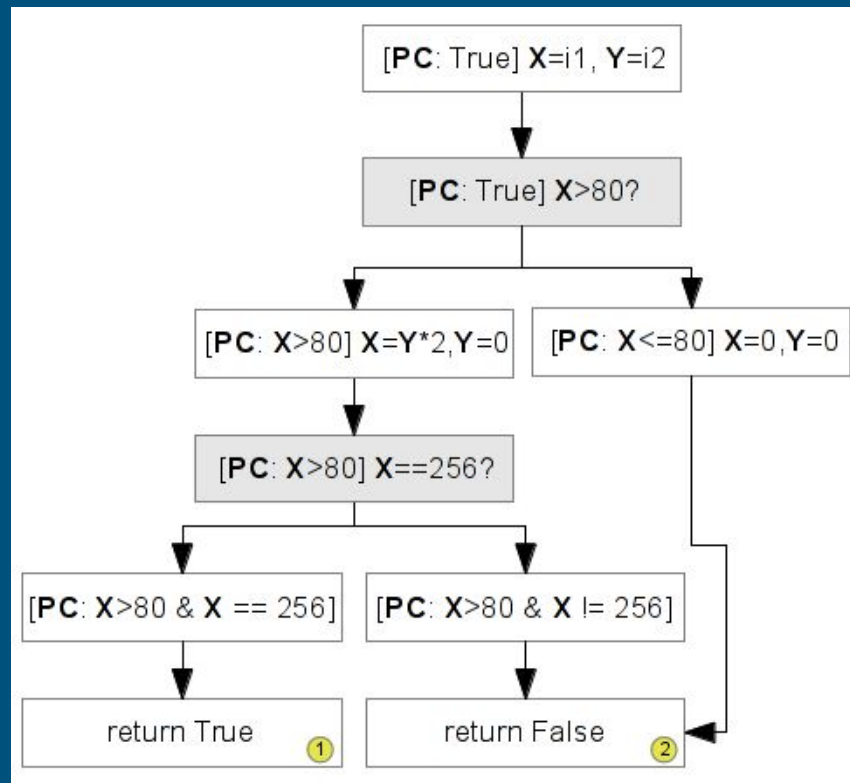
Early Approaches

- Binary Modification
 - 2006: Vasudevan et al. - Cobra
 - Emulate code in blocks like QEMU
 - Remove or rewrite malware instructions that could be used for detection
- Hiding Environmental Artifacts
 - 2007: Willems et al. - CWSandbox
 - In system kernel driver hides environmental artifacts
 - Oberheide later demonstrated several detection techniques against CWSandbox
- State Modification
 - 2009: Kang et al.
 - Builds upon detection work
 - “dynamic state modification” (DSM), modifications to state force malware execution down alternative paths
- Multi-Platform Record and Replay
 - 2012: Yan et al. - V2E
 - Kang et al.’s DSMs are not scalable for multiple anti-analysis checks
 - Don’t mitigate individual occurrences of evasion, make evasion irrelevant because systems are inherently transparent



Path Exploration

- 2007: Moser et al.
 - Looks broadly at code coverage and analyzing trigger-based malware
 - Track when input is used to make control flow decisions, change it to force execution down different paths
- 2008: Brumley et al. - MineSweeper
 - Trigger-based malware focused
 - Represents inputs to potential triggers symbolically, while other code is executed concretely



Hypervisor-based Analysis



- 2008: Dinaburg et al. - Ether
 - Catch system calls and context switches from Xen
 - Despite extensive efforts to make analysis transparent, Pék et al. created nEther and were able to detect Ether
- 2009: Nguyen et al. - MAVMM
 - AMD SVM with custom hypervisor
 - Thompson et al. subsequently demonstrated timing attacks that can be used to detect MAVMM and other hypervisor based systems
- 2014: Lengyel et al. - DRAKVUF
 - Xen-based, instruments code with injected breakpoints

Bare Metal Analysis



- 2011, 2014: Kirat et al. - BareBox & BareCloud
 - BareBox - in-system kernel driver
 - BareCloud - post-run disk forensics
- 2012: Willems et al.
 - Hardware-based branch tracing features
 - Analyzed evasive PDFs
- 2016: Spensky et al. - LO-PHI
 - Instrument physical hardware
 - Capture RAM and disk activity at the hardware level
 - Scriptable user keyboard/mouse interaction with USB-connected Arduinos
- SMM-based analysis: all the transparency benefits of bare metal, while restoring introspection
 - Full access to system memory, protection from modification, high speed, protection from introspection
- 2013 & 2015: Zhang et al. - Spectre, MaIT
 - Spectre: SMM-based analysis, 100x faster than VMM based introspection
 - MaIT - SMM-based *debugging*
- 2016: Leach et al. - Hops
 - SMM memory snapshotting and PCI-based instrumentation

Mitigation - Discussion

- Two broad categories: active and passive mitigation
 - Active: detect-then-mitigate
 - Passive: build inherent transparency
- Passive approaches have been more prevalent
 - Hypervisors, bare metal, etc
- Bare metal is the cutting edge in academic research, but it may not be scalable to industry applications
 - Promising, but not a panacea against any class of attacks other than CPU-based

<https://git.io/vbgEz>

Presentation Outline

1. Introduction
2. Offense - Detecting Analysis Systems
3. Defense - Detecting Malware Evasion
4. Defense - Mitigating Malware Evasion
5. Discussion
6. Conclusion



Discussion

- Offensive Research
 - Reverse Turing Tests
 - Detecting Bare Metal Analysis
- Defensive Research
 - Improving Bare Metal Analysis
 - Heuristic Evasion Detection
 - Passing Reverse Turing Tests
- Research Evaluation
 - Establishing Ground Truth
 - Challenges in research evaluation

Offensive Research

- Reverse Turing Tests
 - Difficult to mitigate against
 - Increasingly relevant as analysis systems become transparent
 - Look to anti-cheating research for online gaming
- Detecting bare metal analysis
 - Still vulnerable to everything except CPU-based attacks
 - Look to detecting analysis instrumentation

Defense - Improving Bare Metal Analysis

- Improving bare metal analysis - efficient, introspection, and stalling mitigation
 - Efficiency
 - 2016: Vadrevu and Perdisci - MAXS - improve efficiency by 50% on average with less than 0.3% information loss in analysis
 - Introspection
 - SMM needs further research
 - Stalling mitigation
 - Difficult to mitigate against with current bare metal systems
 - Performance tracing technologies may provide a direction forward

Defense - Heuristic Evasion Detection

- Can the behaviors involved in evasion before conditional branching occurs be detected heuristically?
- Inspirations
 - Code *fragility* may indicate maliciousness
 - Heuristic detection in enterprise and personal AV/endpoint products
 - Stalling detection techniques
 - Anti-anti-DBI heuristics

Defense - Passing Reverse Turing Tests

- Believably simulating human presence as reverse Turing Tests become more prevalent
- Inspirations:
 - UNVEIL's fake file system creation
 - LARIAT information assurance testbed
 - Biometric spoofing research

Meta - Establishing Ground Truth

- *Unknown-unknowns*: researchers don't know what they don't know
- Human malware analysis is not scalable
- “Bootstrapping” corpora - use previously generated analysis reports as ground truth
 - Problematic: differences in execution environment and time may lead to spurious differences
- Collection in the wild
 - Challenging for *evasive* malware
 - Collection sources may reveal biases

Meta - Challenges in Research Evaluation

- Evaluated works range from evaluating one lab-created malware sample to analyzing millions captured in the wild
- Impossible to empirically compare research, or reproduce results
- 2012: Rossow et al. - evaluated the “methodological rigor and prudence” of 36 papers involving malware experimentation from 2006-2011
 - We re-emphasize all of the author’s points and recommend researchers read their paper closely

<https://git.io/vbgEz>

Conclusion

- Surveyed in paper and backup slides: mobile and web analysis, case studies
- Continual evolution of offense and defense, will to continue into the future
- Cutting edge defenses may not be scalable
- Immense challenges to experimental evaluation and ground truth

ROOTS'17: "A Survey On Automated Dynamic Malware Analysis Evasion and Counter-Evasion: PC, Mobile, and Web"

- Friends who helped us edit: Rolf Rolles, James Kukucka, Aaron Sedlacek
- RPI support: Jeremy Blackthorne and Dr. Greg Hughes
- Program committee and our anonymous reviewers - particularly #4
- Dr. Sergey Bratus and ROOTS PC



@0xAlexei
yener@cs.rpi.edu

Backup Slides

Mobile Analysis - Fingerprinting

- Vulnerable to all of the previously mentioned classes of fingerprinting
- Almost all papers use Android
- Mobile specific:
 - Mobile wear-and-tear artifacts: photos on camera roll, text messages, installed apps, battery life and charge cycles
 - Mobile sensors: multiple radios (cellular, Wi-Fi, Bluetooth, NFC), cameras, microphones, GPS, accelerometers, gyroscopes, thermometers, barometers, proximity sensors, light sensors, magnetic sensors, humidity sensors, air pressure sensors, geomagnetic field sensors, voltmeters
 - Sensor interaction: device vibration can be measured with accelerometer, audio output can be picked up by microphones, CPU strain can increase temperature and strain battery

Mobile Analysis - Defense

- App model means (commodity) mobile malware is fundamentally different than PC malware - it doesn't just "run"
- We could not find any papers on just detecting mobile evasion
- Bare metal analysis:
 - BareDroid - fast system restoration by reverting system drive partitions, monitoring with SELinux
 - A5 - network signature collection with bare metal and virtualized systems
 - Challenges: charging, resetting, introspection

Mobile Analysis - Offensive Research

- Seminal work: Oberheide & Miller's talk on active reconnaissance attacks against Android Bouncer
- Lots of papers on Android analysis detection - almost all just continuing Oberheide & Miller-style attacks
- Advanced fingerprinting research: biometric-based human interaction detection (e.g., variation in touchscreen input) - PIDetector

Web Analysis Systems

- More generally used to look at browser-based exploits, not so much in-browser malicious code
- Two general categories:
 - Instrumented real browsers
 - Look for OS actions indicative of exploitation, e.g., new file or process creation
 - Browser emulators
 - More brittle - emulate the essential parts of a browser - HTTP, HTML, JS, etc with an emulator, look for patterns indicative of exploitation, e.g., a heap spray, or invocation of a function with known bad arguments
- Approaches: submission-based, web crawling, browser plugin analysis

Web Analysis - Fingerprinting

- Use of alternative scripting languages - e.g., ActiveX or VBScript
- Emulators: JS runtime specific bugs, HTML parsing differentials, DOM incompleteness
- HTML5 canvas fingerprinting against *specific* systems
- Shellcode-time detection - detect analysis during shellcode execution
- Referrer checks - only serve content if coming from a specific referrer

Web Analysis - Defense

- Detecting evasion: Kapravelos et al., Revolver
 - Collect web pages *identified* as malicious and benign by existing web analysis systems
 - Compare script ASTs and compare pairwise - if a “benign” script has the same AST as a “malicious” one with the addition of additional control flow nodes, the added nodes are assumed to be evasion checks
- Evasion mitigation
 - Cova et al. - JSAND: force the invocation of JS functions in web malware if not called during regular execution
 - Kolbitsch et al. - Rozzle: multi-path JS forced execution
 - Kim et al. - J-Force: forced path exploration by recording branch outcomes and mutating them. Dynamic creation of DOM elements as needed to prevent crashes on forced paths.

Quantifying Evasion In The Wild

- Extremely difficult to do - lots of “unknown-unknowns” about evasion
 - Likely better done by commercial AV companies seeing millions of samples a day then by academics
- Chen et al. - 2.7% of samples exhibited VMM evasion 2006-2007
- Bayer et al. - 0.03% of samples had a specific anti-Anubis check in 2009
- Chen et al. - anti-analysis became more prevalent 2009-2014, but less prevalent for APTs
- Kapravelos et al. - ~26% evasion rate in web samples

Suggestions for Experimentation

- Establish ground truth
 - Verify analysis results for at least a portion of the malware with a human analyst
- Make multi-execution system similar
 - Minimize differences in environment causing spurious differences in execution
 - Discuss any unavoidable differences
- Be explicit about malware origins
 - Malware corpora may have inherent skews
 - VirusTotal - wild samples caught by defenders, or offensive actors doing testing
 - APTs - hard to catch

Meta - Game Theory Formalizations

- Cat-and-mouse game of analysis system vs. malware
 - Strategy dependent on the “worthiness” of the adversary
 - Save advanced techniques for the most advanced opponent
- Stackelberg games
 - Allocation of analysis resources by analysis system with randomized strategy while malware deploys a purely deterministic evasion strategy

